

The Frontier Overlay — precommodification diagnostic

Goal: Give the dialectic a low-precision **diagnostic overlay** that reads how far each move sits inside vs. outside the latent space — “groove” (precommodified / standard) vs. “frontier” (novel / hard) — and flags the moment a reaching, frontier-ish tension gets *redirected back to known territory*. It is a map the user reads and steers from; it never steers the dialectic itself.

Origin: Venkatesh Rao, “Zero Interest Rate Ideation” (precommodification). Rao’s detection heuristic — *fluent elaboration = precommodified; the model struggling to build on it = frontier* — is the same phenomenon this skill already fights as the “confirms-priors” tendency, stated in a new vocabulary. The skill already implements Rao’s *treatment* (the blind donor sea is his “level-up via combination”; the wiki + tension-queue are his “Zettelkasten + periodic recombination”). What it lacks, and what this feature adds, is a *reading reported back to the user*.

Non-goals / deferred: No auto-steering (overlay only — decided with the user). No per-artifact heat-map over every monk claim (nag risk; the user is a strong detector and reads raw essays fine). No new wiki page type (over-engineering; the reading lives in the control log + inline presentations). A wiki-frontmatter frontier reading is a possible later extension, revisited only after watching this on real runs. This is the concrete first cut of the **open-endedness meter deferred in reference/phase7-recursion.md**.

Design principles (the guardrails — bake these into the wording)

1. **Descriptive, never prescriptive.** It reports terrain; the user decides whether they meant to be on it. It must never nag “this is precommodified, go harder.” On a standard run (user just wants the standard answer) it quietly confirms “groove — as expected” and adds zero friction.
 2. **An ambiguous flag, not a confident score.** Divergence from the expected resolution can mean “frontier” *or* just “the idea was vague/underspecified.” A match can mean “precommodified” *or* just “clear and simple.” The overlay says “*look here,*” not “*this is novel.*” Overclaiming precision would make it worse than nothing.
 3. **An aid to a good detector.** The user detects precommodification and sycophancy readily themselves — this is coverage and attention-direction (catch the collapse on hour three of a long run; let the palette be scanned faster), not a replacement. Low precision and false positives are acceptable because the user filters them at a glance.
 4. **Quiet on standard runs.** The highest-value output — the collapse flag — is silent when there is no frontier reach to collapse *from*. Standard work produces only the cheap “groove” reading.
 5. **Overlay feeds recommendation, never action.** The reading may inform the refinement-loop router’s *recommendation* to the user (which the router already makes), but nothing loops or proceeds on the reading alone. The user remains the only steering authority.
-

What it reads — two layers

Free layer (always on, ~zero cost)

Re-interpret the maturity gate's three existing signals (reference/refinement-loop.md) as a groove/frontier reading:

- Fast saturation / no new cross-edges / no new facts → **groove** (the space is a deep, well-worn well).
- Persistent novelty (new [fit:] cross-edges still appearing, hidden question still moving) → **frontier** (the space is underdetermined; you're on the jagged edge).

Already computed by the router; this layer only *re-labels* it. Per-tension granularity.

Probe layer (fired at the negation and the palette — one cheap blind agent per pass)

A **blind-expectation probe**: an ephemeral agent that sees **only the setup** — the current working question and the poles (the framing of the contradiction) — and nothing else. It predicts the resolution(s) it would *expect* a thoughtful analysis to reach.

- **Input**: working question + the two (or more) committed poles, verbatim. **Blind to**: monk essays, the determinate negation, donor research, the palette, the control log.
- **Task**: “Given this question and these committed positions, what resolution(s) would you expect a competent analysis to land on? Give the 1–3 most likely.” Neutral, no access to the dialectic's actual work.
- **Return**: the 1–3 expected resolutions, concise.
- **Comparison (done by the belief-free orchestrator, which already holds both in context — no extra judge agent)**: each actual negation move / palette candidate is read against the expected set. **Match** → **groove**. **Divergence** → **frontier**.

One probe per refinement pass serves both surfaces, since the setup (question + poles) is constant within a pass.

Decorrelation / firewall. The blind-expectation probe runs *after* the monks and its output goes **only to the orchestrator's frontier reading** — never into a monk brief. It is orchestrator-facing like donor/tension/synthesis material. It is ephemeral (recorded in the reading + ledger, not a wiki page).

The collapse detector (the highest-value output)

The redirect-to-known-territory the user fears has a signature:

- **Reach present**: the monks reached (hard-won / heavily-qualified / donor-dependent argumentation, or new [fit:] cross-edges appeared in the 4.6 decomposition), **AND**
- **Groove landing**: the determinate negation's result — or a palette candidate — matches the blind-expected resolution.

When both hold, flag it: “**🚩 reached on X — negation/candidate pulled back to the groove Y.**” Detecting “the monks reached” is a qualitative orchestrator judgment (was the argument fluent

and complete, or did it strain and invent?) — an honest flag, not a metric. On a standard run there is no reach, so no collapse flag appears.

Where it surfaces

1. Negation — a companion block (primary)

Where the user spends the most reading time, and the collapse *site* (Aufhebung is exactly where a reaching tension is either elevated into something new or “resolved” into a groove). When the orchestrator presents the determinate negation, it appends a **short frontier-reading companion block** *after* the negation (keeps the negation itself clean to read; puts the diagnostic in one scannable place the user can ignore at will). Contents:

- The groove/frontier reading (free layer) in one line.
- The blind-expected resolution(s) vs. where the negation actually landed.
- Any collapse flag(s).
- 3–4 lines total. Points at moves *by name*; does not re-tag the negation text inline.

2. Palette — one-line flag per candidate (secondary)

At the Phase 5 decision point, each S/J/G/F/U candidate gets a one-line groove/frontier flag (from the same blind expectation). Lets the user scan “which of these five is actually off the worn surface.” Naturally, S (the fluent, obvious synthesis) will often read groove; U/F often read frontier — the flag makes that visible per-run rather than assumed.

3. Control log — a frontier-ledger (cross-pass tracker)

Reuses the hidden-question-ledger machinery. One line per refinement pass, so redirect *over time* is visible (a snapshot can’t see that a frontier idea from pass 1 was flattened by pass 3). Schema (added to the control-log section of [reference/dialectic-wiki.md](#)):

frontier-ledger:

```
- pass N: reading=<groove|frontier|mixed>; expected=<blind expectation, one line>; actual=<where negation/candidate landed>; collapse=<none | "reached on X, pulled to groove Y">
```

Files touched

- **reference/frontier-overlay.md (NEW)**. The central doc: the blind-expectation probe contract (input/blindness/task/return), the orchestrator comparison, the free layer, the collapse detector, the companion-block format, the frontier-ledger schema, and the five guardrails. Everything else points here.
- **reference/phase4-determinate-negation.md + the relevant negation stage (phase4-stage-c-decomposition.md / phase4-stage-d-criteria.md)**. After the determinate negation is produced and before it is presented: spawn the blind-expectation probe, write the companion block, add the frontier-ledger line. Add a completion-gate item (“frontier reading attached to the negation; blind-expectation probe was blind to monk/negation material”).

- **reference/phase5-sublation.md.** Add the one-line groove/frontier flag per candidate; gate item.
 - **reference/refinement-loop.md.** The router’s control-log update writes the frontier-ledger line each pass; the router’s *recommendation* may cite the reading (esp. a collapse) as diagnostic evidence — without acting on it. Add gate item.
 - **reference/dialectic-wiki.md.** Add the frontier-ledger schema to the per-round control-log section; note the blind-expectation probe as orchestrator-facing/ephemeral (firewall).
 - **SKILL.md.** Brief mention: the Frontier Overlay, what it reads, where it surfaces, and the “descriptive not prescriptive / overlay not operator” framing. Add to the phase map / summaries where the negation and palette are described.
-

Self-review notes

- **Placeholder scan:** none — probe contract, comparison owner, collapse logic, ledger schema, and companion-block contents are all concrete.
- **Consistency:** firewall treatment of the blind-expectation probe matches the existing donor/tension/synthesis orchestrator-facing rule; “overlay feeds recommendation not action” matches the router’s existing “diagnose and recommend; user decides” contract; “no new page type” matches the deferred-extension stance.
- **Scope:** single, focused feature (a diagnostic overlay); one new reference doc + edits to five existing files. Fits one implementation plan.
- **Ambiguity:** “the monks reached” is deliberately a qualitative orchestrator judgment, consistent with guardrail 2 (flag, not metric) — called out rather than left implicit.